

An intro to
Map-Reduce

Romaric Duvignau

Associate Professor

Computer & Network Systems, CSE

Slides from Ahmed Ali-Eldin (slightly updated)



CHALMERS
UNIVERSITY OF TECHNOLOGY



Distributed Computing and Systems
Chalmers university of technology

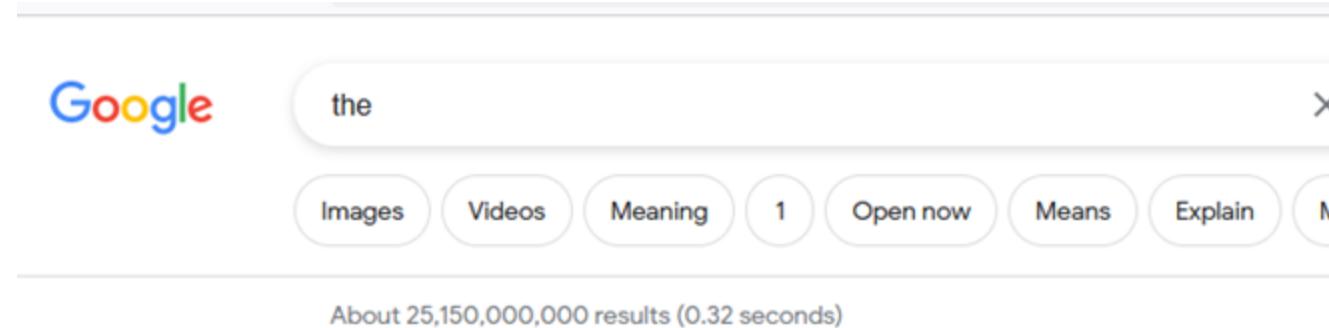
Example: Counting Words

- Count the most popular terms on the Internet:
- Sounds easy
- Why not this?

```
CountWords(internet) :  
    hash_map<string, int> word_count;  
    for each page in internet:  
        for each word in page:  
            word_count[word] += 1
```

Motivation: Large-Scale Data Processing

- I searched for the word “the”



Motivation: Large-Scale Data Processing

- All the web = ~40 billion* web pages x 2.5MB** =
from 100 Petabytes (1 PB = 1000 TB) up to 100 ZB (total information)
- The largest hard disk money can buy today is a 125 TB, but that is very rare, so we go for the more standard SSD 30 TB
- ~3000 hard drives just to store the small web terabytes
- One computer can read up to 1 Gbps ~ 2.63 Pb/month from disk
 - ~38 months to read the hard drives sequentially from a single machine with no failures
 - (but Google gave me an answer in 0.32s!)

*<https://www.worldwidewebsize.com/>

**<https://www.pingdom.com/blog/webpages-are-getting-larger-every-year-and-heres-why-it-matters/>

REMINDER: Don't use AI to do your 5th grade Maths!

Google

1 gbps in one month

Mode IA Tous Images Produits Vidéos Actualités Vidéos courtes Plus Outils

◆ Aperçu IA

1 Gbps is equal to approximately **2.63 terabits per month**, or about 3.24×10^{14} bytes per month. To convert this, first calculate the number of seconds in a month (around 2,629,744) and multiply that by 1 Gbps (1 billion bits per second). [🔗](#)



Calculations

- **Seconds in a month:**

$$30.44 \text{ days} \times 24 \text{ hours/day} \times 60 \text{ minutes/hour} \times 60 \text{ seconds/minute} \approx 2,629,744 \text{ seconds}$$

- **Gigabits per month:**

$$1 \text{ Gbps} \times 2,629,744 \text{ seconds} = 2,629,744,000,000 \text{ bits/month}$$

- **Terabits per month:**

$$2,629,744,000,000 \text{ bits} \div 1000 = 2,629,744 \text{ gigabits} = 2,629.744 \text{ terabits}$$

- **Bytes per month (Base 10):**

$$1 \text{ Gbps} \times 2,629,744 \text{ seconds} \div 8 \text{ bits/byte} = 328,718,000,000 \text{ bytes/month} = 3.287 \times$$



Large-Scale Data Processing

1. Problems:

1. Time
2. Errors / Problems

2. Solution:

1. Distribute processing on many machines

Typical things companies do

1. Indexing algorithms

- a) Which web pages contain a certain word

2. Link Reversal

- a) Which pages link to this page

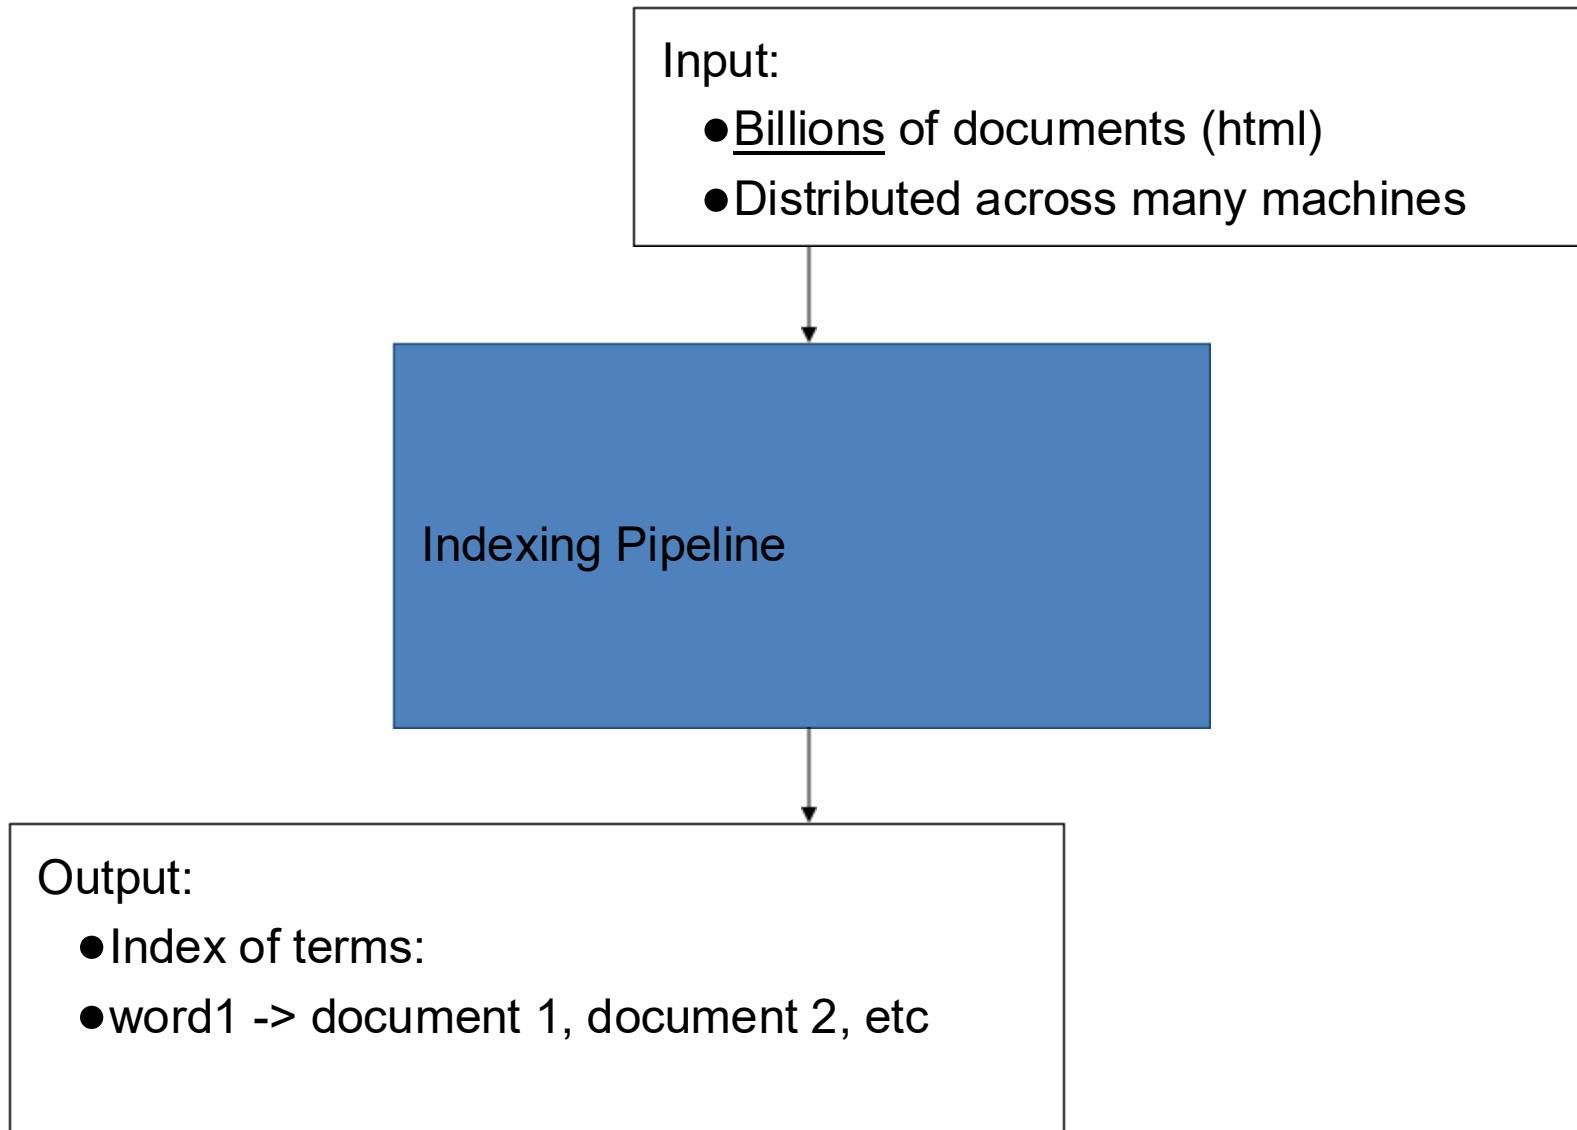
- For example

- “<https://www.cse.chalmers.se/~duvignau/>”

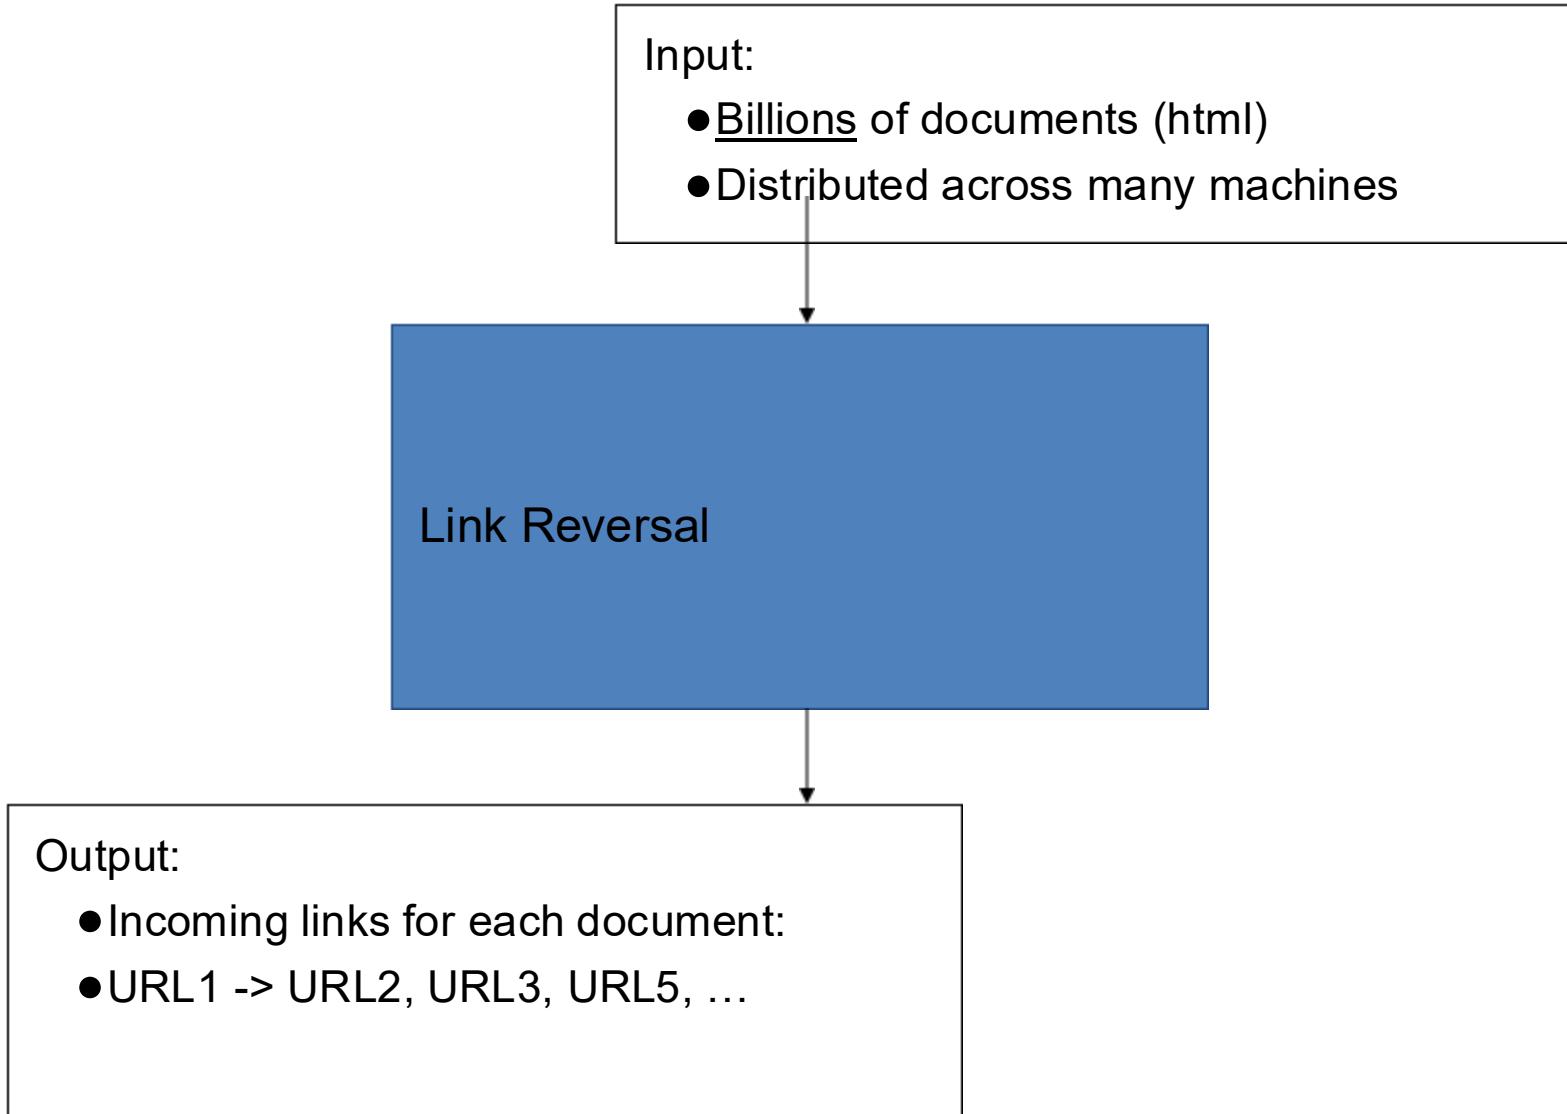
- ← “Chalmers, Google scholar, linkedin, Facebook, twitter, etc”

3. Sorting large data alphabetically

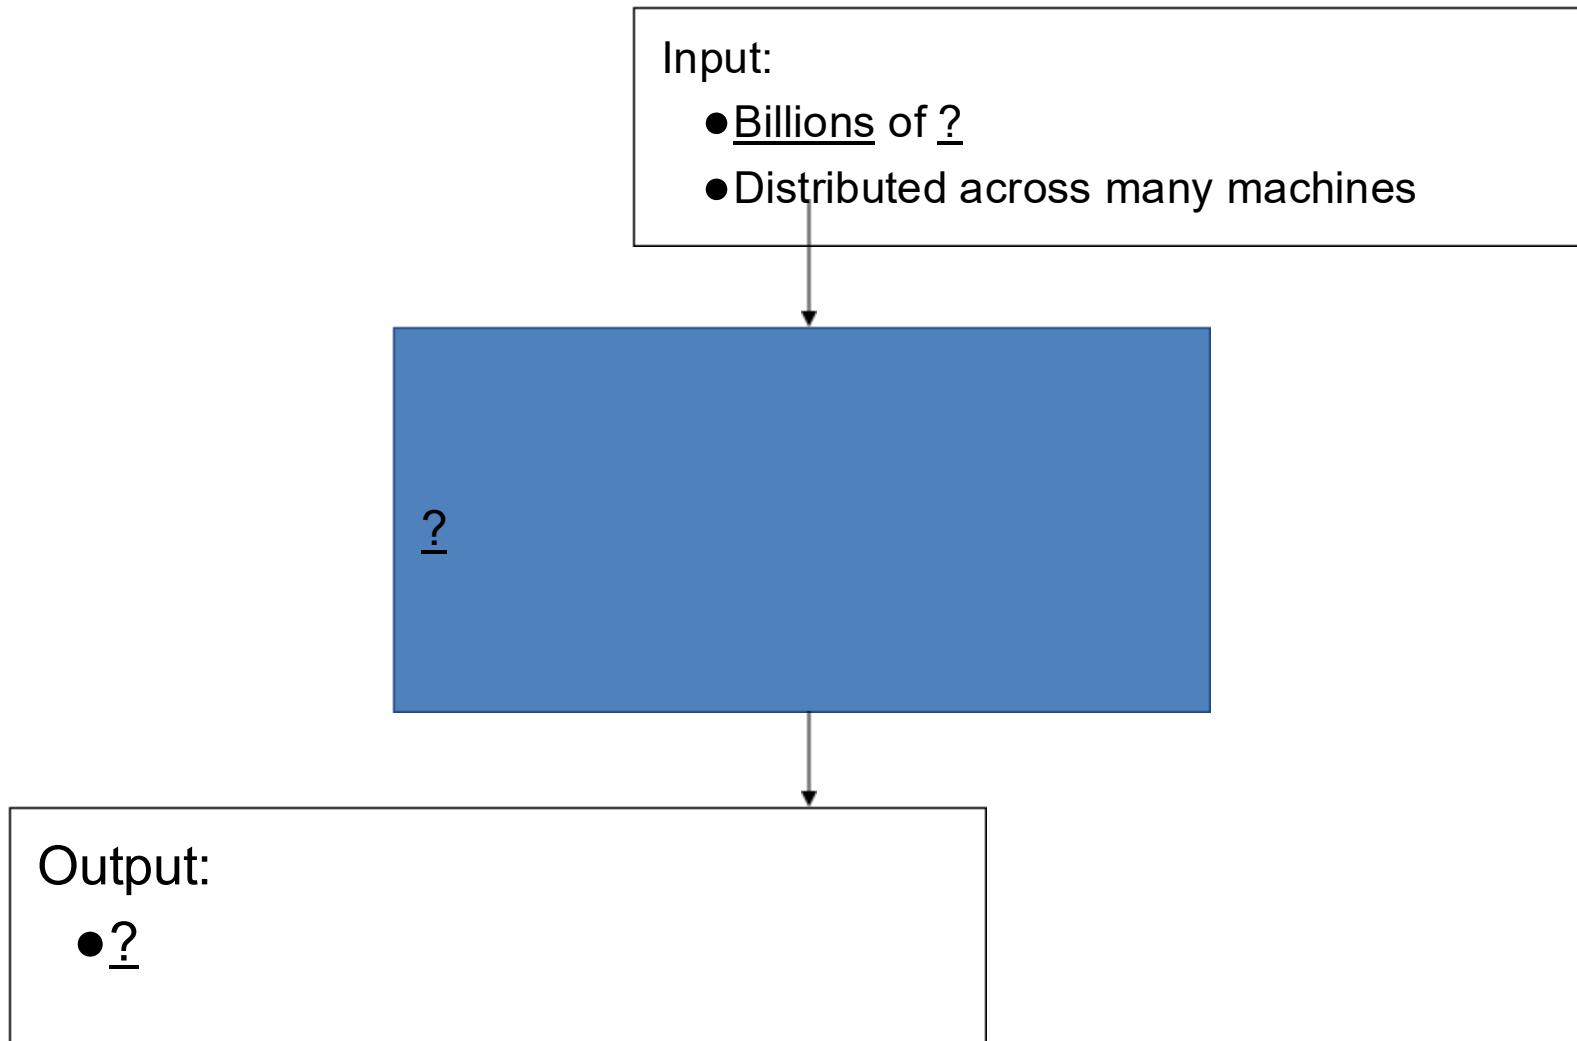
Indexing Algorithm



Link Reversal



Any Algorithm



Before MapReduce

- Each team needed to understand distributed computing
- Extra code
- Extra complexity
-

Algorithm

- Requirements:
 - Has to be scalable (terabytes of data)
 - Has to be fast
 - Has to be fault-tolerant (data & machines)
 - Has to be easy to implement and flexible
- The solution:
 - **MapReduce: a distributed algorithm**

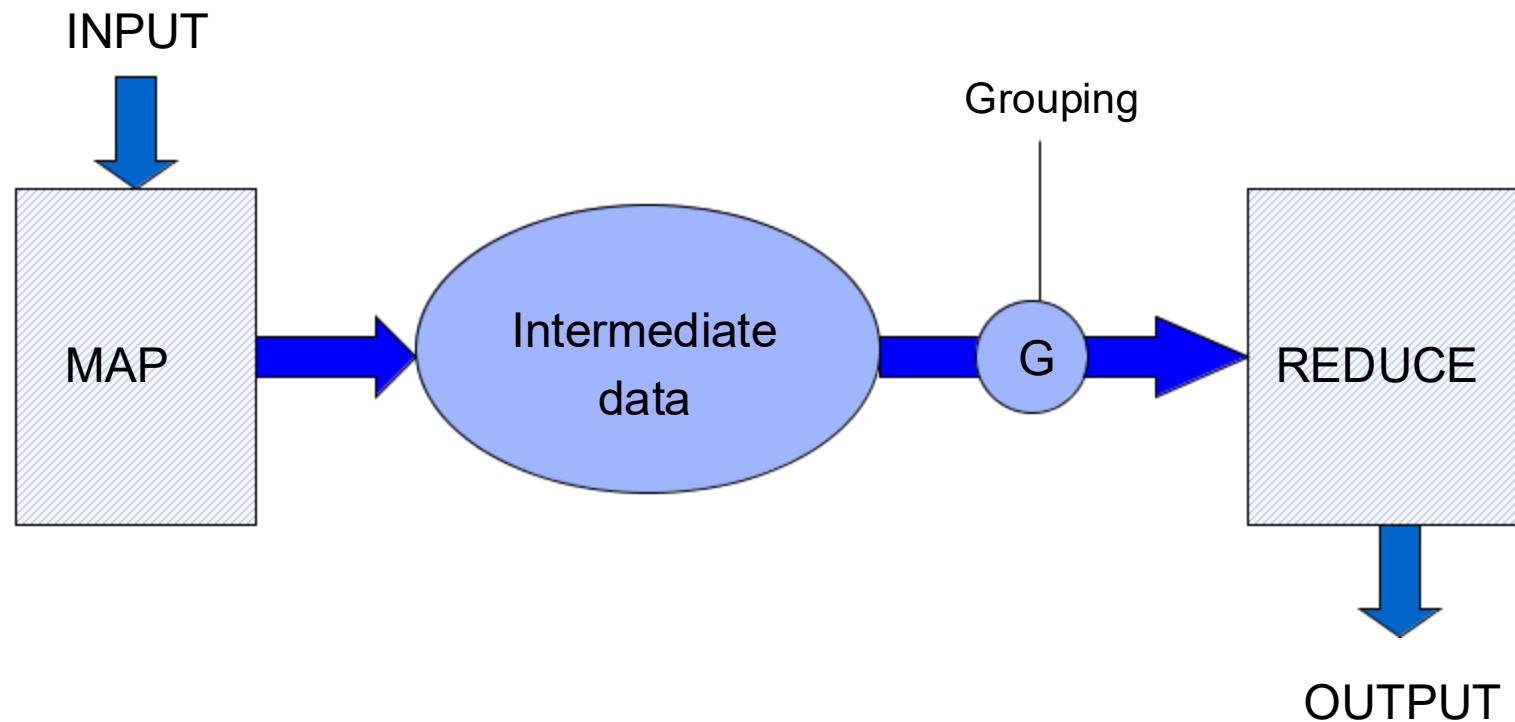
MapReduce: Introduction

- “*MapReduce is a programming model and an associated implementation for processing and generating large data sets.*”
- Can run on thousands of machines
- Can process terabytes of information (input / output)
- It is fault-tolerant
- It is easy to implement
- It is flexible

MapReduce Model

- Programming model inspired by functional language primitives, e.g., LISP
 - 1. Read a lot of data
 - 2. **Map**: extract something you care about from each record
 - 3. Shuffle and Sort
 - 4. **Reduce**: aggregate, summarize, filter, or transform
 - 5. Write the results
-
- Outline stays the same, map and reduce change to fit the problem

MapReduce: Logical Diagram



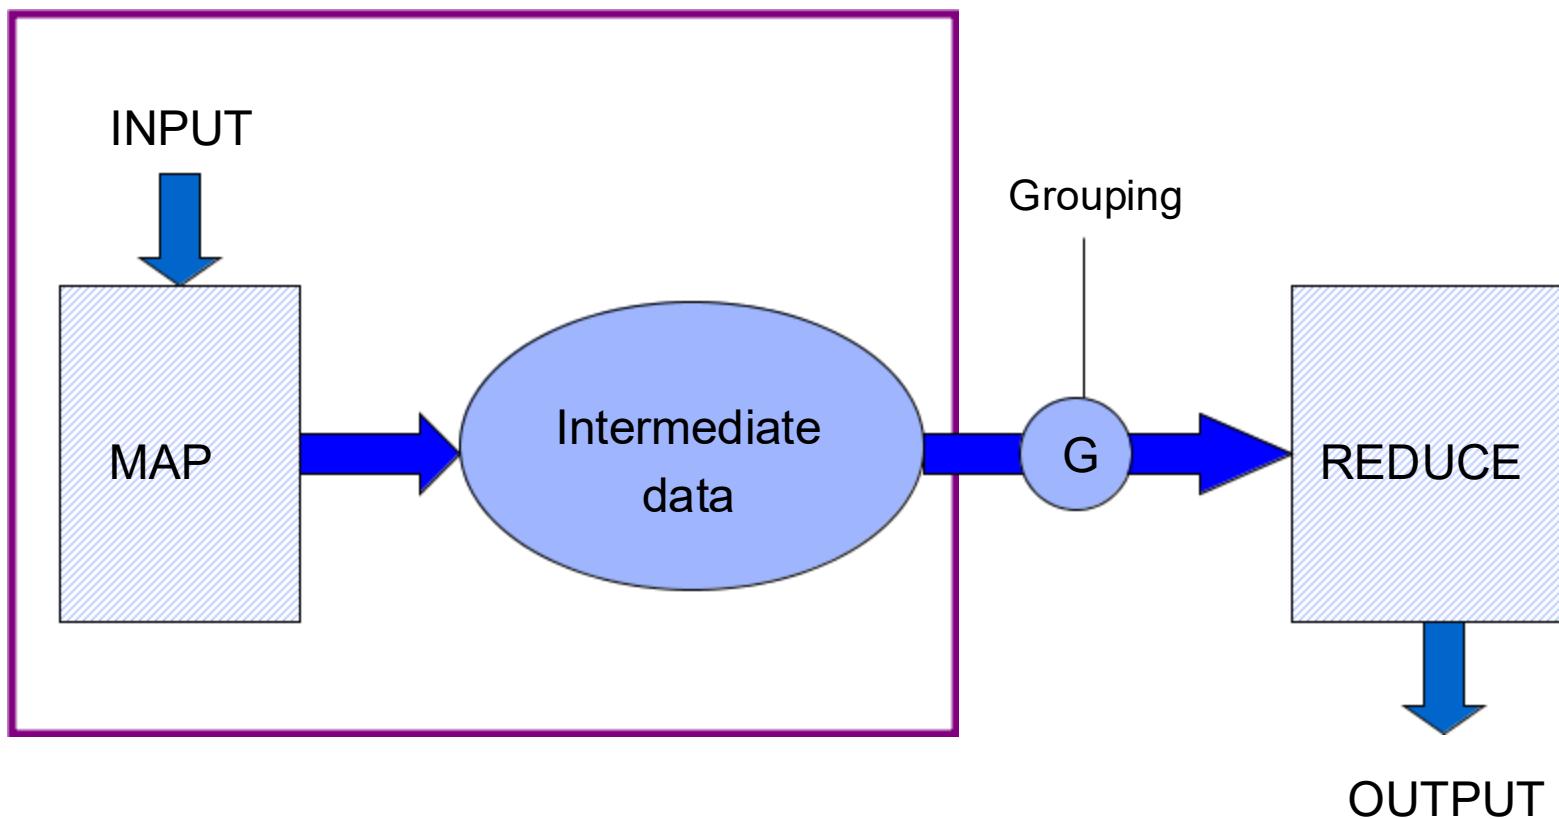
Programming Model

- Input & Output: each a set of key/value pairs

Programmer specifies two functions:

- **map (in_key, in_value) → list(out_key, intermediate_value)**
 - Processes input key/value pair
 - Produces set of intermediate pairs
- **GROUPING**
 - Done by the system
- **reduce (out_key, list(intermediate_value)) → list(out_value)**
 - Combines all intermediate values for a particular key
 - Produces a set of merged output values (usually just one)

MapReduce: Logical Diagram



MapReduce: Map Phase

- First, the **MAP** function:
- Programmer supplies this function
 - `map (in_key, in_value)`
→ `list(out_key, intermediate_value)`
- Input = any pair (e.g., URL, document)
- Output = list of pairs (key/value)

MapReduce: Map Example

- Example: word counting
 - <URL, document text> → Output: list(<word, count>)
- E.g. for one document
 - <URL1, "Chalmers university is the best university">
 - Emit → list(<Chalmers, 1>, <university, 1>, <is, 1>, <the, 1>, <best, 1>, <university, 1>)
- For another document
 - <URL2, "UNIX is simple and elegant">
 - Emit → list(<unix, 1>, <is, 1>, <simple, 1> . . .)

MapReduce: Map Example

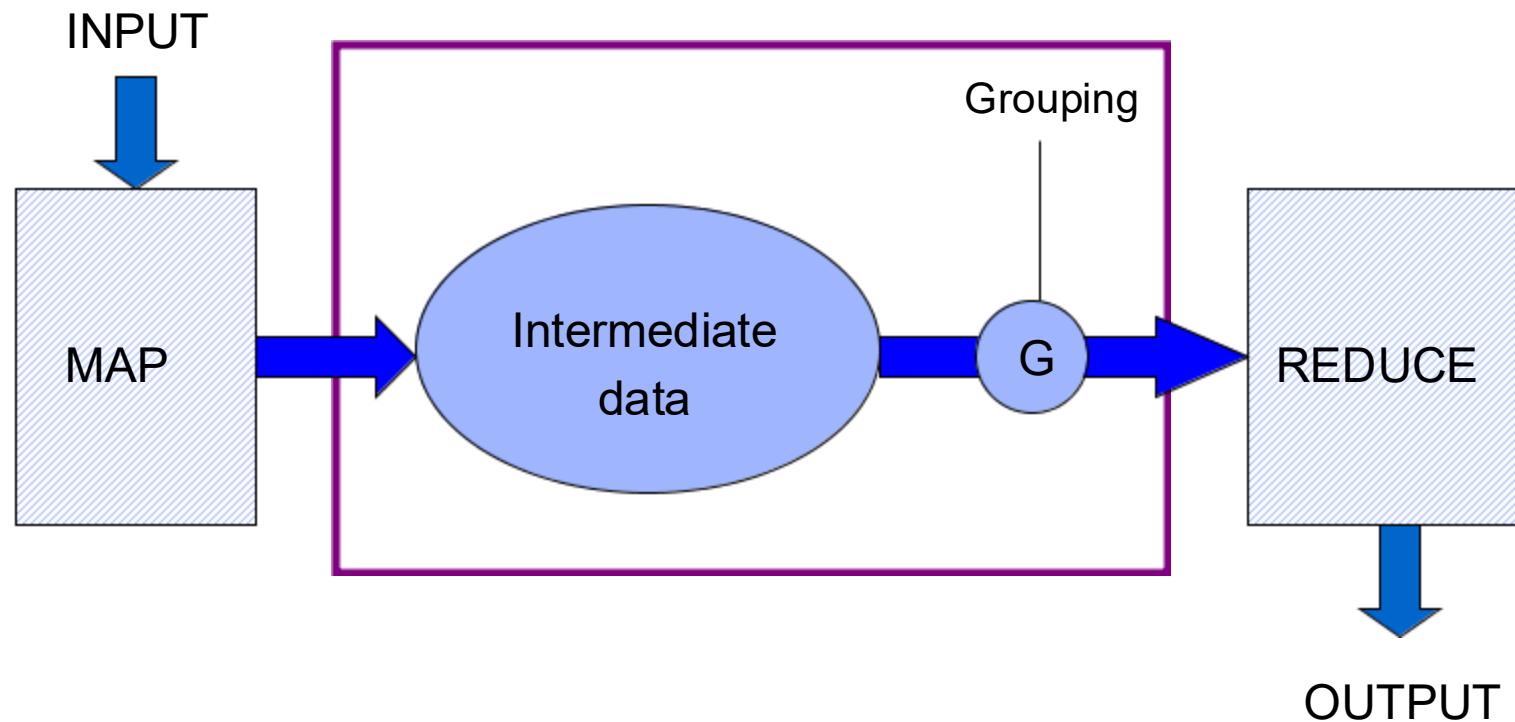
- Example: word counting

- <URL, document text>
- Output: list(<word, count>)

- Code:

```
Map(string input_key, string input_value):  
    // key: document URL  
    // value: document text  
    for each word in input_value:  
        OutputIntermediate(word, "1");
```

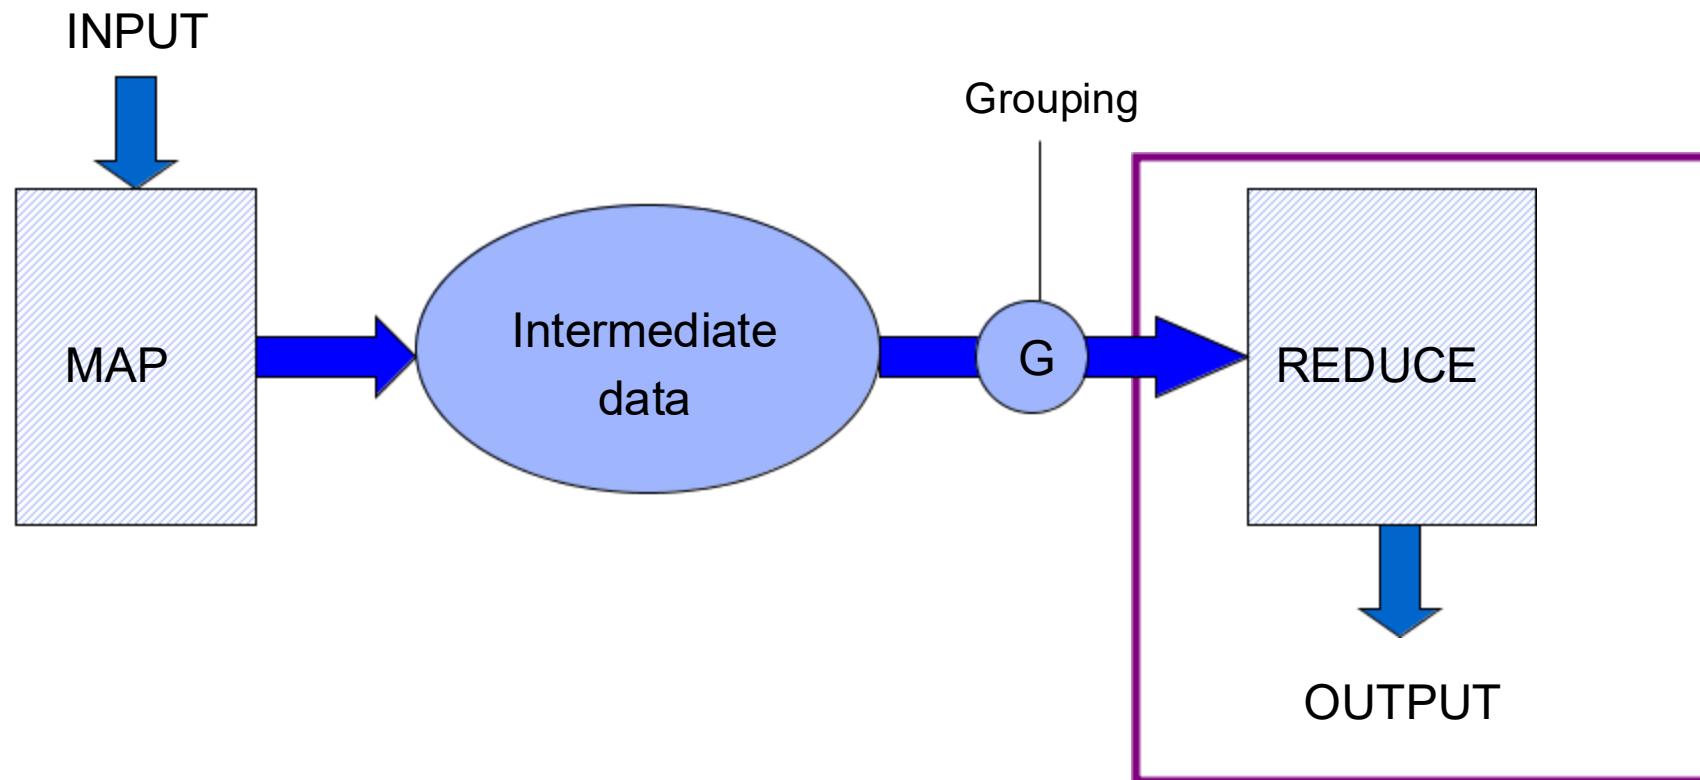
MapReduce: Logical Diagram



MapReduce: Grouping

- The **GROUPING** (or **shuffling**)
- This happens automatically (by the library)
- Put all values for same key together
- For our example:
 - <Chalmers, list(1)>
 - <university, list(1, 1)>
 - <the, list(1)>
 - <is, list(1, 1)>
 - <best, list(1)>
 - ...

MapReduce: Logical Diagram



MapReduce: Reduce

- Last phase is **REDUCE**
- Takes the output from **grouping**:
- `reduce (out_key, list(intermediate_value)) → list(out_value)`

MapReduce: Reduce Example

- Example (word-counting):

- <university, list(1, 1)> • <university, 2>

```
Reduce(string key, List values):  
    // key: word, same for input and output  
    // values: list of counts  
    int sum = 0;  
    for each v in values:  
        sum += v;  
    Output(sum);
```

MapReduce: Word Count

- Example's summary: Want to count how many times each word occurs:

- **MAP:**

- Input: <**DocumentId**, **Full Text**>
 - For each time a word occurs:
 - Output: <**Word**, 1>

- **REDUCE:**

- Input: <**Word**, list(**1, 1, ...**)>
 - Sum over each word:
 - Output <**Word**, **Count**>

MapReduce: Grep (search)

- Want to search terabytes of documents for where a word occurs:

- **MAP:**

- Input: <**DocumentId**, **Full Text**>
 - For each line where the word occurs:
 - Output: <**DocumentId**, **LineNumber**>

- **REDUCE:**

- Input: <**DocumentId**, list(**LineNumber**, ...)>
 - Do nothing (*Identity* function, i.e. output = input)

MapReduce: Link Reversal

- Want to build a reverse-link index:

- **MAP:**

- Input: <**source_URL**, **Full Text**>
 - For each outgoing link :
 - Output: <**destination_URL**, **source_URL**>

- **REDUCE:**

- Input: <**destination_URL**, list(**source_URL**, ...)>
 - Join source_URLs
 - Output: for each URL, a list of all other URLs that link to it.

MapReduce Usage

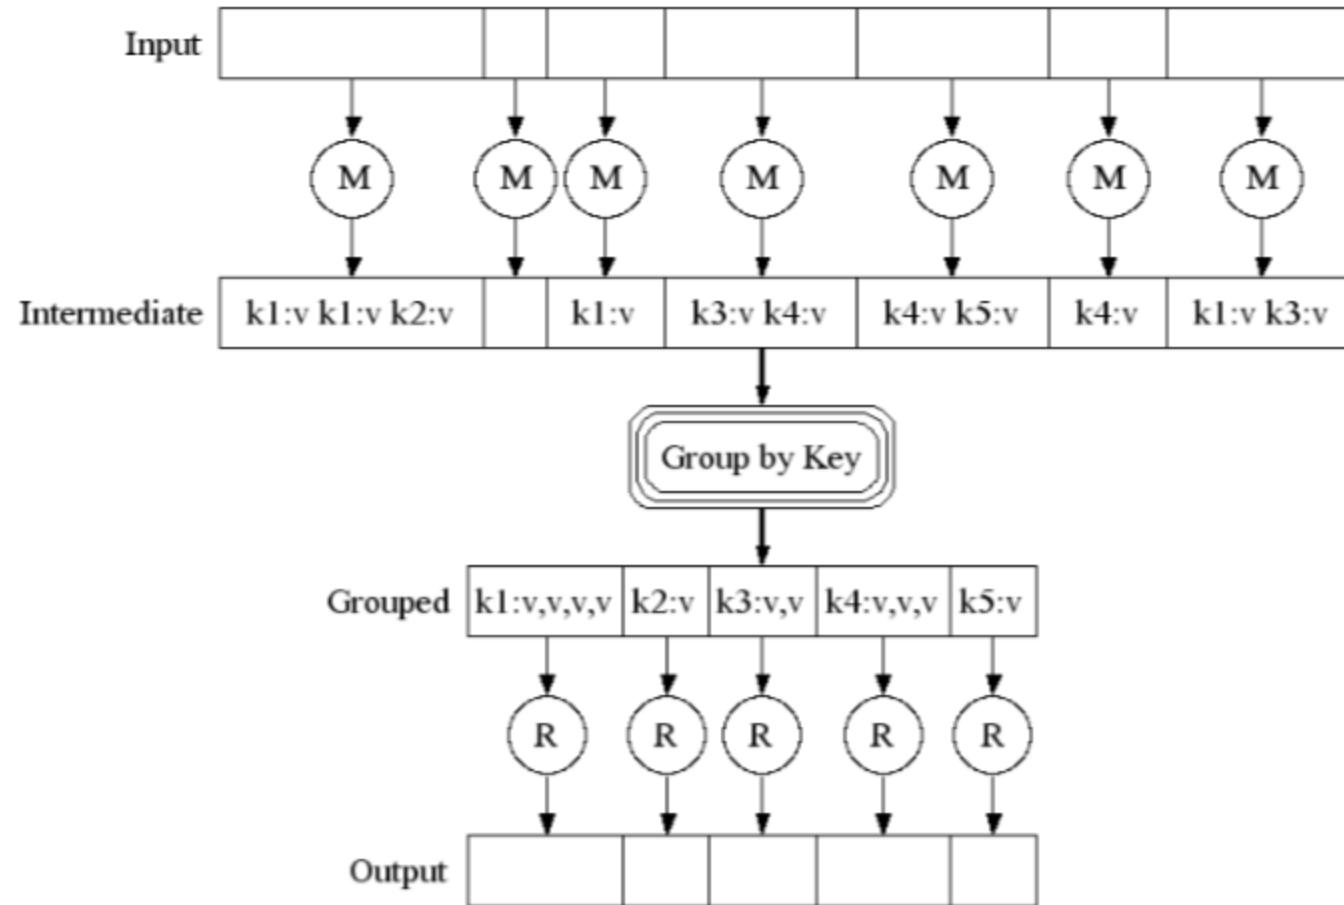
- Index construction
- Web access log stats, e.g., search trends (Google Zeitgeist)
- Distributed grep
- Distributed sort
- Most popular terms on the Internet
- Document clustering (news, products, etc.)
- Find all the pages that link to each document (link reversal)
- Machine learning
- Statistical machine translation
- ...

Why complicate things?

MapReduce: Parallelism

- Each Map call is independent (parallelize)
 - Run **M** different **Map tasks**
(each takes a chunk of input; make use of locality!)
- Each Reduce call is independent (parallelize)
 - Run **R** different **Reduce tasks**
(each produces a chunk of output; there are R output files.)

MapReduce: Parallelism



Framework

- User:
 - **Map** function (used in Map Phase)
 - **Reduce** function (used in Reduce Phase)
 - Options (Map tasks, Reduce tasks, Workers, Input, Output)
- Framework:
 - Execute map and reduce phase (distributedly)
 - Collect/move data:
 - input
 - intermediate
 - output data
 - Handle faults
 - Dynamic load-balancing and scheduling

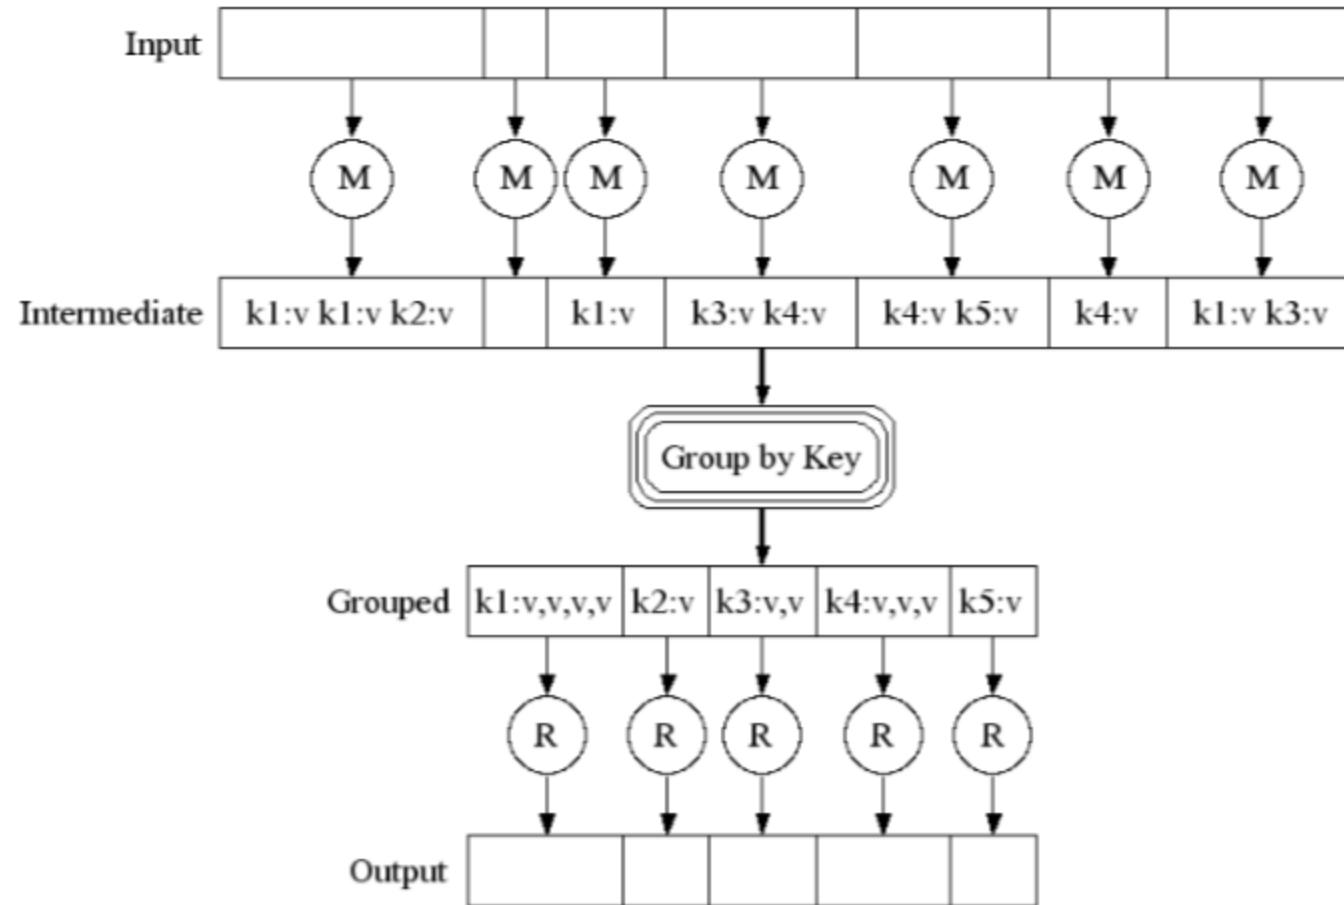
MapReduce: Scheduling

- One coordinator node C , many workers W_1, W_2, \dots
 - Input data split into M map tasks (16-64MB in size)
 - Reduce phase partitioned into R reduce tasks
 - Tasks are assigned to workers dynamically
 - Often: $M=200,000$; $R=4,000$; workers=2,000
- Coordinator assigns each map task to a free worker
 - Considers locality of data to worker when assigning task
 - Worker reads task input (often from local disk! or *Distributed File System!*)
 - Worker produces R local files (spill files) containing intermediate k/v pairs
- Coordinator assigns each reduce task to a free worker
 - Worker reads intermediate k/v pairs from map workers
 - Worker sorts & applies user's *Reduce* operation to produce the output

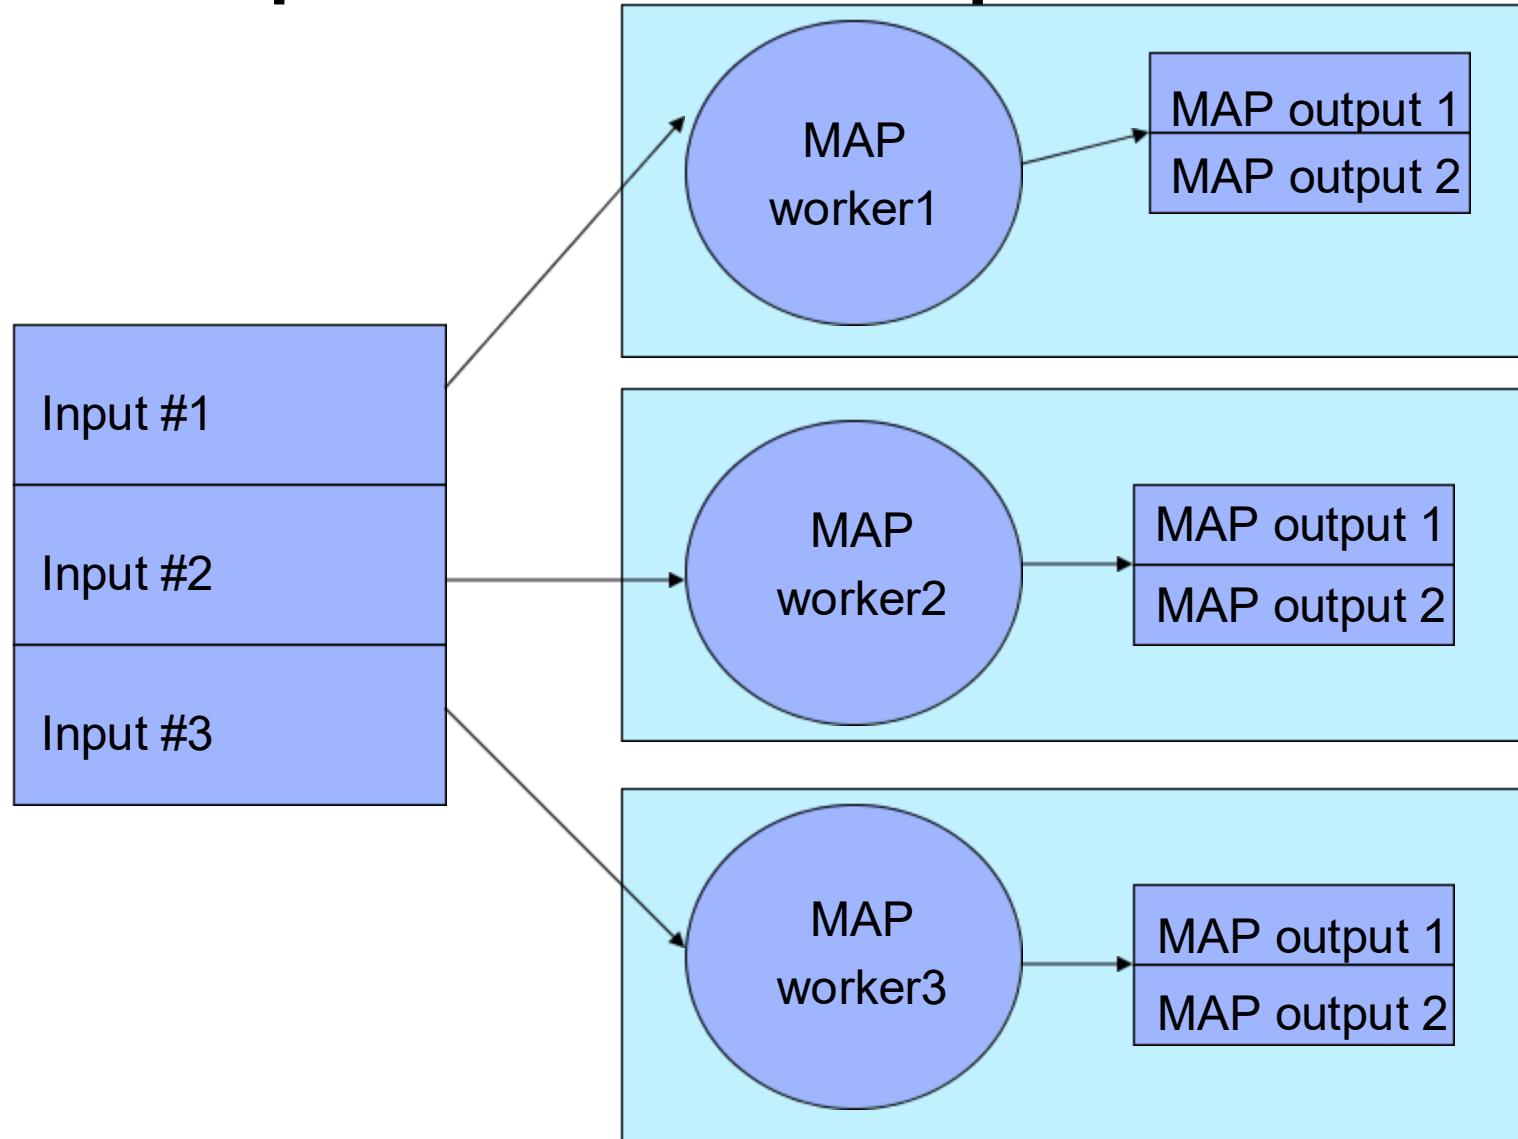
MapReduce: The “Coordinator”

- One machine is the “Coordinator” = Control Plane
- Controls other machines (workers), start or stop MAP/REDUCE
- Keeps track of progress of each worker
- Stores location of input, output of each worker
- There are MAP workers and REDUCE workers

MapReduce: Parallelism

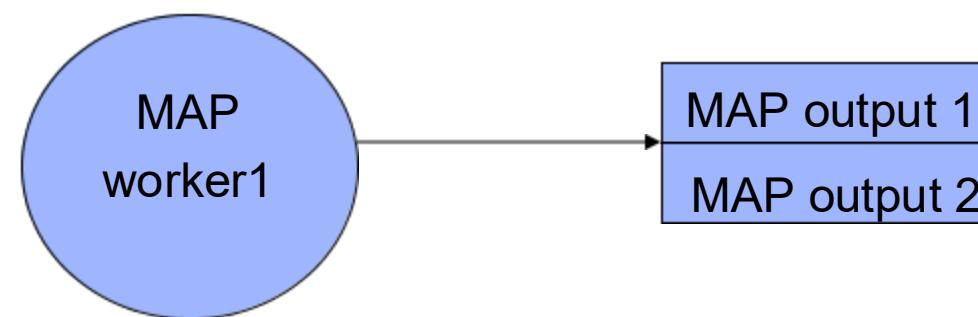


MapReduce: Map Parallelism

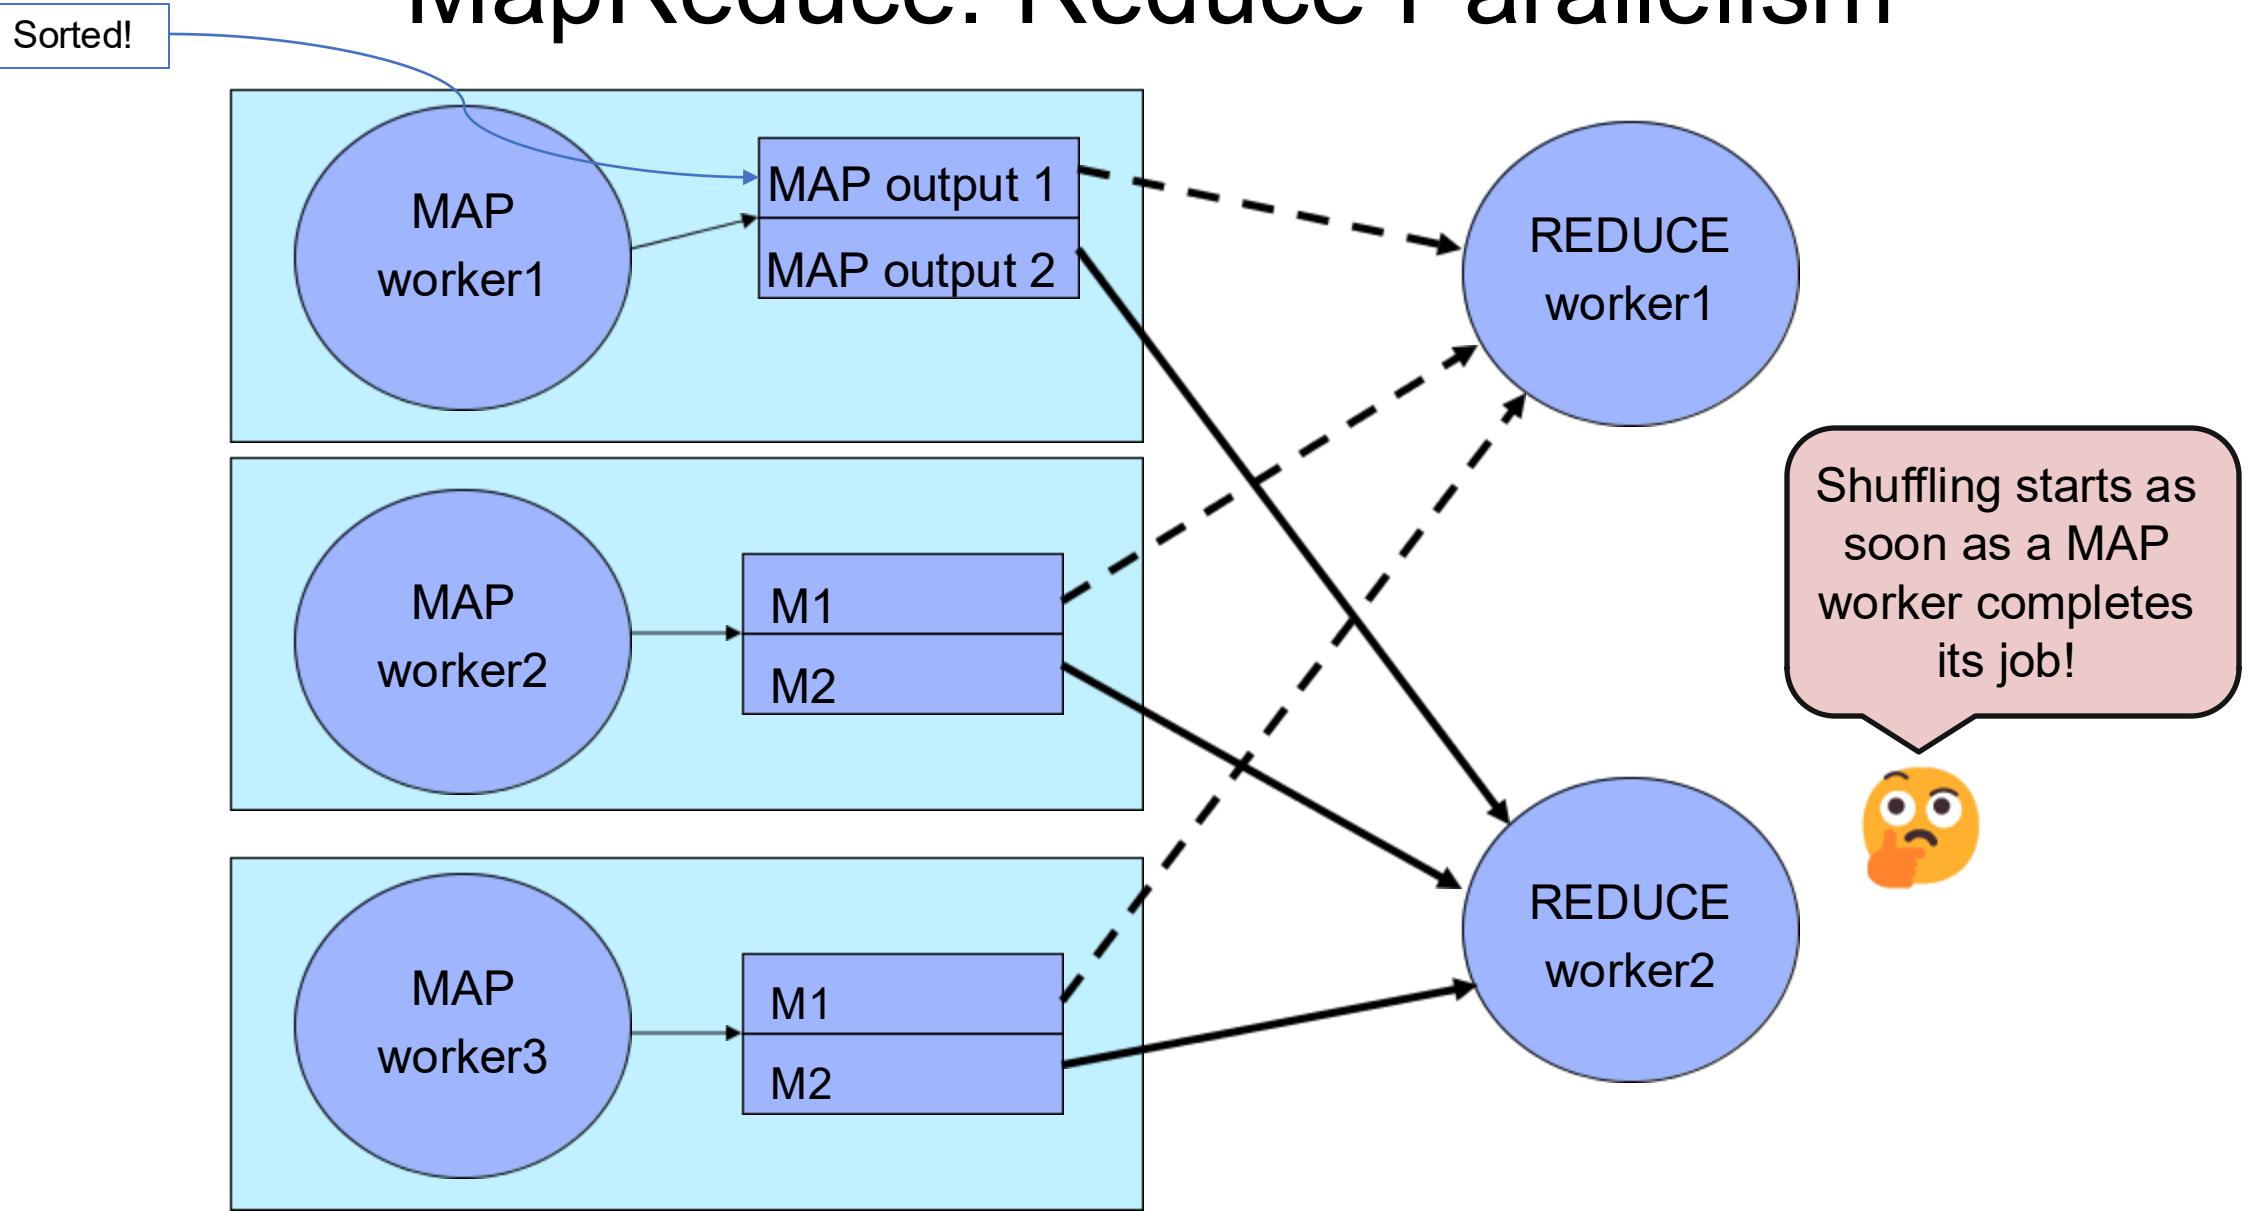


MapReduce: Map Parallelism

- MAP output is split for REDUCERs by key
- “Bins” for keys
 - E.g. keys starting $(a - m)$ → **output1** and $(n - z)$ → **output2**
- (Usually use $\text{hash}(\text{key}) \bmod (\# \text{ of reducers})$ for more random partition)

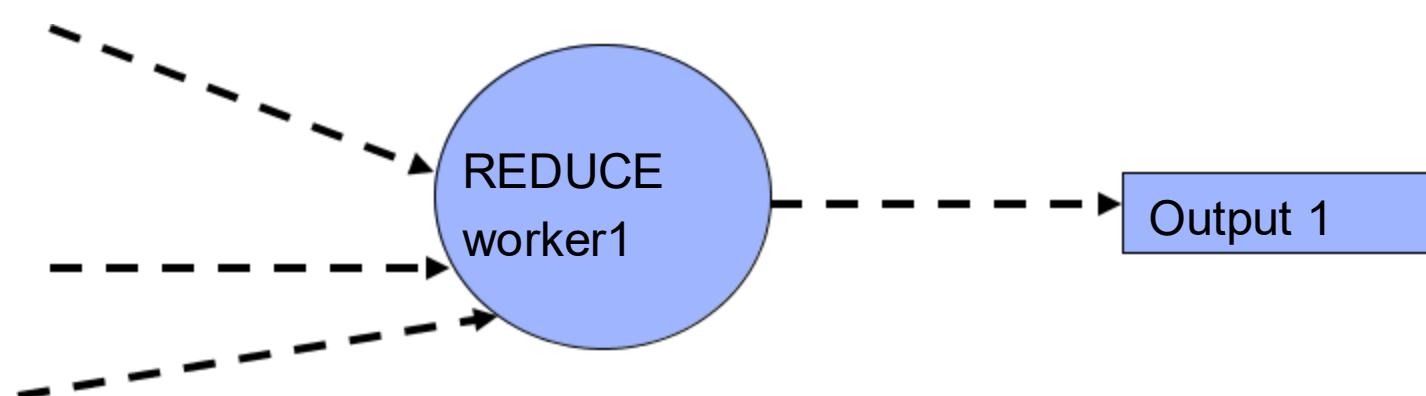


MapReduce: Reduce Parallelism



MapReduce: Reduce Parallelism

- Each REDUCE worker reads one section of input from all MAP workers
 - Performs GROUPing
 - Performs reducing
 - Stores output



MapReduce: Fault Tolerance

- More machines = more chance of failure
- Probability of one machine failing → p (small)
- If each machine fails independently, chance of any one out of k machine failing → $1 - (1 - p)^k$ (larger than p)
- If $p = 0.01$, probability of one failure:
 - $1 \text{ machine} \cdot 0.01 = 1\%$ a machine will fail
 - $2 \text{ machines} \cdot 0.0199 = 2\%$ a machine will fail
 - $10 \text{ machines} \cdot 0.096 = 10\%$ a machine will fail
 - $100 \text{ machines} \cdot 0.63 = 2/3$ chance a machine will fail
 - $1000 \text{ machines} \cdot \mathbf{0.9999} = \mathbf{99.99\% \ chance!}$

MapReduce: Fault Tolerance

“If one machine fails, the program should not fail”

“If many machines fail, the program should not fail”

MapReduce: Fault Tolerance

- Worker fails (happens often):
 - Coordinator checks workers
 - If worker responds as failed, or does not respond:
 - Coordinator restarts that worker on different machine
- Bad/buggy input:
 - Coordinator keeps track of bad input
 - Asks workers to skip bad input
- Coordinator fails (happens rarely):
 - Can abort operation
 - Or, can restart the coordinator on a different machine

MapReduce: Optimizations

- No REDUCE can start until all MAPs are complete (bottleneck)
 - At least the second part (after shuffling)
- Sometimes one machine is really slow (bad RAM, corrupted disk, etc)
 - Coordinator notices slow machines (called **stagglers**)
 - Executes input on another machine
 - Uses output of machine that finishes first
- Other optimizations
 - Locality optimization
 - Intermediate data compression ("combiner" pre network transfer)

Task Granularity and Pipelining

- Fine granularity tasks: many more map tasks than machines
 - e.g., 200,000 map/5000 reduce tasks w/ 2000 machines
- Minimizes time for fault recovery
- Can pipeline grouping with map execution
- Better dynamic load balancing

MapReduce: Beyond Google

- Hadoop: Open-source implementation of MapReduce
 - Hadoop v1: reimplementations of Google's MapReduce
 - Hadoop v2: separation of the Control Plane (Coordinator, resource management, etc) to YARN
- MapReduce implementations in C++, Java, Python, Perl, Ruby, Erlang
- ...
- Most companies that run distributed algorithms have an implementation of MapReduce (Facebook, Nokia, MySpace, ...)

Extra: Let us take a look at how to code with MR

- We will show code for Hadoop streaming running Python code for word count.
 - Hadoop was built for Java
 - Hadoop streaming is the way you write code in other languages:
<https://hadoop.apache.org/docs/r1.2.1/streaming.html>
- Example from here:
 - <https://www.geeksforgeeks.org/hadoop-streaming-using-python-word-count-problem/>

mapper.py

Python3

```
#!/usr/bin/env python

# import sys because we need to read and write data to STDIN and STDOUT
import sys

# reading entire line from STDIN (standard input)
for line in sys.stdin:
    # to remove leading and trailing whitespace
    line = line.strip()
    # split the line into words
    words = line.split()

    # we are looping over the words array and printing the word
    # with the count of 1 to the STDOUT
    for word in words:
        # write the results to STDOUT (standard output);
        # what we output here will be the input for the
        # Reduce step, i.e. the input for reducer.py
        print '%s\t%s' % (word, 1)
```

reducer.py

Python3

```
#!/usr/bin/env python

from operator import itemgetter
import sys

current_word = None
current_count = 0
word = None

# read the entire line from STDIN
for line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()
    # splitting the data on the basis of tab we have provided in mapper.py
    word, count = line.split('\t', 1)
    # convert count (currently a string) to int
    try:
        count = int(count)
    except ValueError:
        # count was not a number, so silently
        # ignore/discard this line
        continue

    # this IF-switch only works because Hadoop sorts map output
    # by key (here: word) before it is passed to the reducer
    if current_word == word:
        current_count += count
    else:
        if current_word:
            # write result to STDOUT
            print '%s\t%s' % (current_word, current_count)
        current_count = count
        current_word = word

    # do not forget to output the last word if needed!
if current_word == word:
    print '%s\t%s' % (current_word, current_count)
```

Some comments

- Here is another example with optimizations using iterators and generators in Python from: <https://www.michael-noll.com/tutorials/writing-an-hadoop-mapreduce-program-in-python/>

mapper.py

```
#!/usr/bin/env python
"""A more advanced Mapper, using Python iterators and generators."""

import sys

def read_input(file):
    for line in file:
        # split the line into words
        yield line.split()

def main(separator='\t'):
    # input comes from STDIN (standard input)
    data = read_input(sys.stdin)
    for words in data:
        # write the results to STDOUT (standard output);
        # what we output here will be the input for the
        # Reduce step, i.e. the input for reducer.py
        #
        # tab-delimited; the trivial word count is 1
        for word in words:
            print '%s%s%d' % (word, separator, 1)

if __name__ == "__main__":
    main()
```

reducer.py

```
#!/usr/bin/env python
"""A more advanced Reducer, using Python iterators and generators."""

from itertools import groupby
from operator import itemgetter
import sys

def read_mapper_output(file, separator='\t'):
    for line in file:
        yield line.rstrip().split(separator, 1)

def main(separator='\t'):
    # input comes from STDIN (standard input)
    data = read_mapper_output(sys.stdin, separator=separator)
    # groupby groups multiple word-count pairs by word,
    # and creates an iterator that returns consecutive keys and their group:
    #   current_word - string containing a word (the key)
    #   group - iterator yielding all ["<current_word>", "<count>"] items
    for current_word, group in groupby(data, itemgetter(0)):
        try:
            total_count = sum(int(count) for current_word, count in group)
            print "%s%s%d" % (current_word, separator, total_count)
        except ValueError:
            # count was not a number, so silently discard this item
            pass

if __name__ == "__main__":
    main()
```

Java example?

```
WordCount.java
1. package org.myorg;
2.
3. import java.io.IOException;
4. import java.util.*;
5.
6. import org.apache.hadoop.fs.Path;
7. import org.apache.hadoop.conf.*;
8. import org.apache.hadoop.io.*;
9. import org.apache.hadoop.mapred.*;
10. import org.apache.hadoop.util.*;
11.
12. public class WordCount {
13.
14.     public static class Map extends MapReduceBase implements Mapper<LongWritable, Text, Text, IntWritable> {
15.         private final static IntWritable one = new IntWritable(1);
16.         private Text word = new Text();
17.
18.         public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output, Reporter reporter) throws IOException {
19.             String line = value.toString();
20.             StringTokenizer tokenizer = new StringTokenizer(line);
21.             while (tokenizer.hasMoreTokens()) {
22.                 word.set(tokenizer.nextToken());
23.                 output.collect(word, one);
24.             }
25.         }
26.     }
27.
28.     public static class Reduce extends MapReduceBase implements Reducer<Text, IntWritable, Text, IntWritable> {
29.         public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text, IntWritable> output, Reporter reporter) throws IOException {
30.             int sum = 0;
31.             while (values.hasNext()) {
32.                 sum += values.next().get();
33.             }
34.             output.collect(key, new IntWritable(sum));
35.         }
36.     }
37.
38.     public static void main(String[] args) throws Exception {
39.         JobConf conf = new JobConf(WordCount.class);
40.         conf.setJobName("wordcount");
41.
42.         conf.setOutputKeyClass(Text.class);
43.         conf.setOutputValueClass(IntWritable.class);
44.
45.         conf.setMapperClass(Map.class);
46.         conf.setCombinerClass(Reduce.class);
47.         conf.setReducerClass(Reduce.class);
48.
49.         conf.setInputFormat(TextInputFormat.class);
50.         conf.setOutputFormat(TextOutputFormat.class);
51.
52.         FileInputFormat.setInputPaths(conf, new Path(args[0]));
53.         FileOutputFormat.setOutputPath(conf, new Path(args[1]));
54.
55.         JobClient.runJob(conf);
56.     }
57. }
58.
59.
```

C++ code example

```
#include "hadoop/Pipes.hh"
#include "hadoop/TemplateFactory.hh"
#include "hadoop/StringUtils.hh"

class WordCountMap: public HadoopPipes::Mapper {
public:
    WordCountMap(HadoopPipes::TaskContext& context){}
    void map(HadoopPipes::MapContext& context) {
        std::vector<std::string> words =
            HadoopUtils::splitString(context.getInputValue(), " ");
        for(unsigned int i=0; i < words.size(); ++i) {
            context.emit(words[i], "1");
        }
    }
};

class WordCountReduce: public HadoopPipes::Reducer {
public:
    WordCountReduce(HadoopPipes::TaskContext& context){}
    void reduce(HadoopPipes::ReduceContext& context) {
        int sum = 0;
        while (context.nextValue()) {
            sum += HadoopUtils::toInt(context.getInputValue());
        }
        context.emit(context.getInputKey(), HadoopUtils::toString(sum));
    }
};

int main(int argc, char *argv[]) {
    return HadoopPipes::runTask(HadoopPipes::TemplateFactory<WordCountMap,
                                WordCountReduce>());
}
```

Can I play with it?

YES!

Please do using EMR which is available for you on the AWS lab we have. Documentations:

<https://docs.aws.amazon.com/emr/latest/ManagementGuide/emr-what-is-emr.html>

I like this, what should I read?

- Spark:

https://www.usenix.org/legacy/event/hotcloud10/tech/full_papers/Zaharia.pdf

- You can get an online free playground for spark here

- Databricks: <https://docs.databricks.com/en/getting-started/community-edition.html>

- EMR:

<https://docs.aws.amazon.com/emr/latest/ReleaseGuide/emr-spark.html>

I like this, what should I read?

- Flink (coming from KTH!): <https://www.diva-portal.org/smash/get/diva2:1059537/FULLTEXT01.pdf>
 - You can get an online free playground for spark here
 - Flink:
<https://docs.aws.amazon.com/emr/latest/ReleaseGuide/emr-flink.html>